



# Interaction Terms & Hierarchical Linear Models

Fitting Models to Data, Not Data to Models

Model Fitting Series - With Applications in R

---

Jesse Ghashti

October 21, 2025

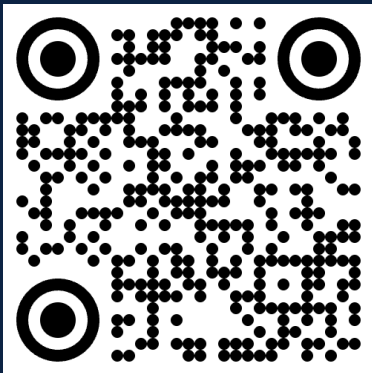
Centre for Scholarly Communication

The University of British Columbia | Okanagan Campus | Syilx Okanagan Nation Territory

## Install Required Packages

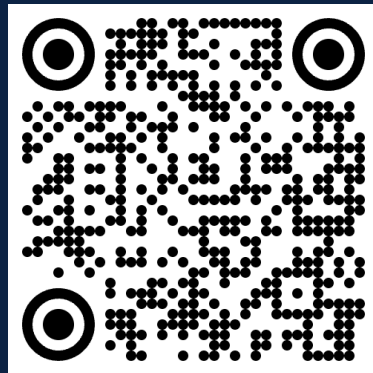
```
packages <- c("dplyr", "mgcv", "ggplot2",  
             "gratia", "janitor")  
toInstall <- packages[!(packages %in%  
                        installed.packages()[,"Package"])]  
if(length(toInstall)) install.packages(toInstall)  
  
library('dplyr')    # for data wrangling  
library('mgcv')     # for modeling  
library('ggplot2')  # for fancy plots  
library('janitor')  # for data cleaning  
library('gratia')   # ggplot-based graphics
```

Session	Topic	Date/Time
1	Simple Linear Regression	Oct 7, 9:00 AM
2	Fitting Linear Models in R	Oct 8, 10:30 AM
3	Multiple Linear Regression in R	Oct 16, 4:00 PM
4	<b>Interaction Terms &amp; Hierarchical Linear Models</b>	<b>Oct 21, 11:00 AM</b>
5	Generalized Linear Models	Oct 23, 4:00 PM
6	Generalized Additive Models (GAMs)	Oct 28, 11:00 AM
7	Interpreting & Predicting from GAMs	Oct 29, 10:30 AM
8	Hierarchical GAMs	Nov 4, 12:00 PM
9	Penalized Models	Nov 18, 11:00 AM
10	Survival Models	Nov 25, 11:00 AM
11	Nonparametric Models	Dec 2, 11:00 AM



←  
New to R? Check  
out the Fundamen-  
tals of R series!

GitHub code and  
slides for today's  
workshop (and pre-  
vious workshops)



Alternatively, code/slides available at the bottom of  
<https://csc-ubc-okanagan.github.io/workshops/>



### Key Concepts

- Extended to multiple linear regression:  $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \epsilon$
- Used DAGs to visualize causal relationships
- Learned about multicollinearity and VIF
- Interpreted partial effects (holding other variables constant)
- Saw that more predictors  $\neq$  better model (use  $R^2_{adj}$ )

# Today: Interaction Terms & Hierarchical Models



## Today we will...

- Learn when effects depend on context (interactions)
- Distinguish fixed effects vs random effects
- Build hierarchical linear models (mixed models)
- Work with `ChickWeight` data: growth curves by diet
- Visualize and interpret interaction effects

## Today we require...

```
library('dplyr')    # for data wrangling
library('mgcv')     # for modeling
library('ggplot2')  # for fancy plots
library('gratia')   # for ggplot-based model graphics
library('janitor')  # for cleaning names
theme_set(theme_classic(base_size = 15))
```

# What Are Interaction Terms?



## When Effects Depend on Context

**Without interaction:**  $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \epsilon$

- Effect of  $x_1$  is always  $\beta_1$ , regardless of  $x_2$

**With interaction:**  $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2 + \epsilon$

- Effect of  $x_1$  depends on  $x_2$ :  $\frac{\partial Y}{\partial x_1} = \beta_1 + \beta_3 x_2$


**Real-world example:** The effect of study time on grades might depend on teaching method

# Exploring ChickWeight Data



```
cw <- janitor::clean_names(ChickWeight) %>%  
  as_tibble()  
# Plot growth by diet  
ggplot(cw, aes(time, weight)) +  
  facet_wrap(~ diet) +  
  geom_line(aes(group = chick)) +  
  geom_smooth(method = 'lm',  
             formula = y ~ x,  
             color = 'darkorange')
```

Each diet shows different growth patterns — perfect for interactions!





# Model o: Ignoring Diet Differences I



```
# Model that only accounts for time
m_cw_0 <- gam(formula = weight ~ time,
               family = gaussian(),
               data = cw,
               method = 'ML')
```

```
summary(m_cw_0)
```

This model assumes all chicks grow at the same rate!

```
weight ~ time
```

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	27.4674	3.0365	9.046	<2e-16 ***
time	8.8030	0.2397	36.725	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.7 Deviance explained = 70.1%

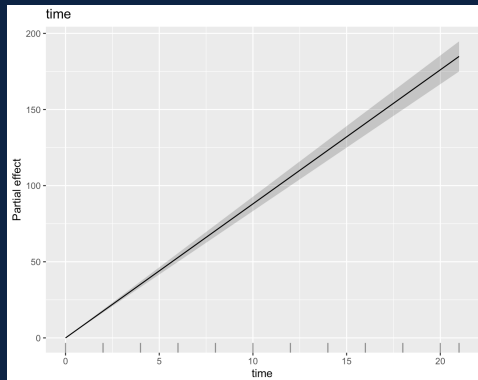
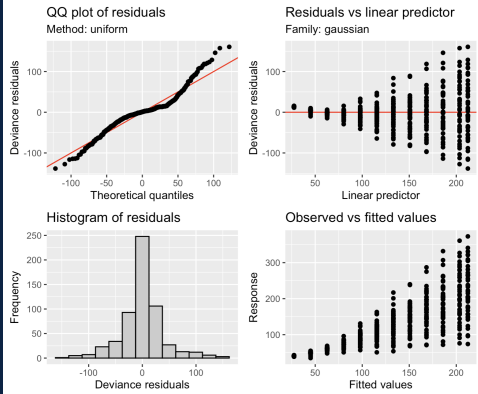
-ML = 2935.4 Scale est. = 1514.3 n = 578

# Model o: Ignoring Diet Differences II



```
appraise(m_cw_0)
```

```
draw(m_cw_0, parametric = TRUE)
```



## Fixed Effects

- Independent categories
- We care about specific levels
- Examples:
  - Treatment vs Control
  - Male vs Female
  - Diet 1, 2, 3, 4
- Each gets its own parameter

## Random Effects

- Levels from a population
- We care about variation
- Examples:
  - Individual subjects
  - Schools in a district
  - Chicks in experiment
- Share information (pooling)

Random effects allow "borrowing strength" across levels

# Model 1: Adding Diet & Chick Effects I



```
m_cw_1 <- gam(  
  formula = weight ~  
    time +           # temporal trend  
    diet +           # fixed: diet type  
    s(chick, bs = 're'), # random: chicks  
  family = gaussian(),  
  data = cw,  
  method = 'ML')  
summary(m_cw_1)  
# Note changes in R-sq, Dev. expl.
```

Formula:  
weight ~ time + diet + s(chick, bs = "re")

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	11.2311	5.5858	2.011	0.04487 *
time	8.7175	0.1755	49.672	< 2e-16 ***
diet2	16.2193	9.0916	1.784	0.07499 .
diet3	36.5527	9.0916	4.020	6.65e-05 ***
diet4	30.0255	9.0984	3.300	0.00103 **

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(chick)	39.91	46	7.641	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.842    Deviance explained = 85.4%  
-ML = 2802.6    Scale est. = 800.06    n = 578

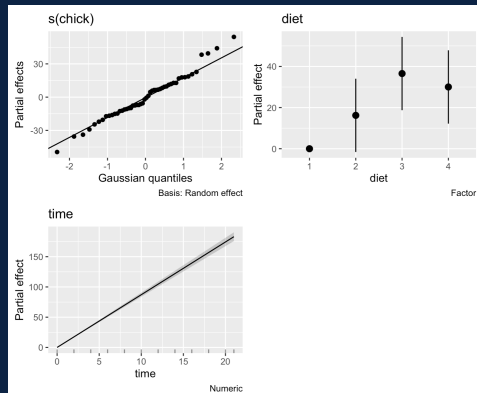
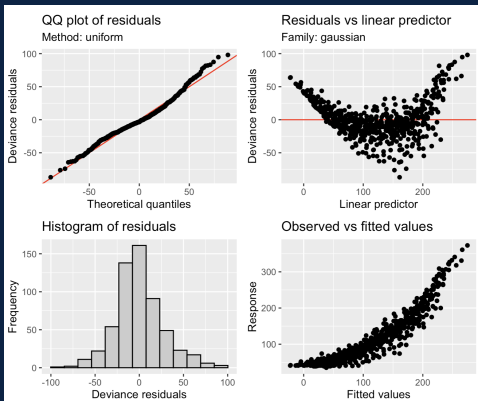
`s(chick, bs = 're')` creates random intercepts — each chick can start at a different weight

# Model 1: Adding Diet & Chick Effects II



```
appraise(m_cw_1)
```

```
draw(m_cw_1, parametric = TRUE)
```



## Model 2: Diet Affects Growth Rate I



```
m_cw_2 <- gam(formula = weight ~  
  time +  
  diet +  
  time:diet +  
  s(chick, bs = 're'),  
  family = gaussian(),  
  data = cw,  
  method = 'ML')  
summary(m_cw_2)
```

Formula:

weight ~ time + diet + time:diet + s(chick, bs = "re")

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	31.5081	5.9358	5.308	1.63e-07 ***
time	6.7130	0.2584	25.982	< 2e-16 ***
diet2	-2.8745	10.2342	-0.281	0.779
diet3	-13.2577	10.2342	-1.295	0.196
diet4	-0.3983	10.2430	-0.039	0.969
time:diet2	1.8961	0.4285	4.425	1.17e-05 ***
time:diet3	4.7099	0.4285	10.992	< 2e-16 ***
time:diet4	2.9495	0.4340	6.795	2.92e-11 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(chick)	41.15	46	9.879	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.873 Deviance explained = 88.3%

-ML = 2744 Scale est. = 643.72 n = 578

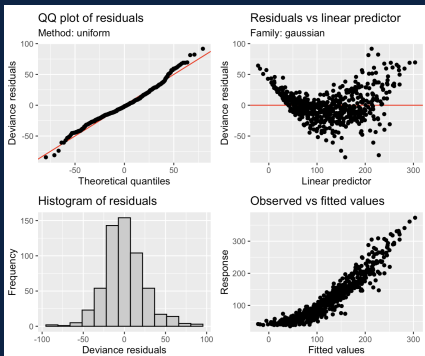
time:diet means each diet has its own slope — some diets lead to faster growth?

# Model 2: Diet Affects Growth Rate II

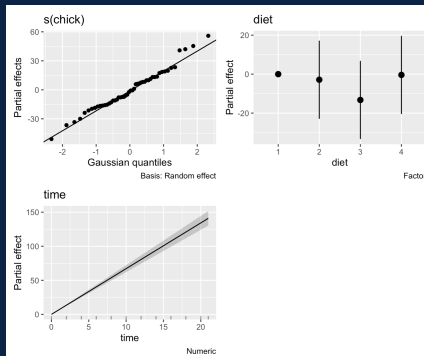


Interpretability of `draw()` might be a little abstract since no interaction terms are plotted...

```
appraise(m_cw_2)
```




```
draw(m_cw_2, parametric = TRUE)
```



# Visualizing Interaction Effects



```
# plot predictions for "average" chick
m_cw_2 %>%
  data_slice(time = evenly(time),
             diet = evenly(diet),
             chick = 'new') %>%
  fitted_values(object = m_cw_2,
               data = .) %>%
  ggplot() +
  facet_wrap(~ diet) +
  geom_ribbon(aes(time,
                ymin = .lower_ci,
                ymax = .upper_ci),
            alpha = 0.3) +
  geom_line(aes(time, .fitted),
           color = 'darkorange')
```





## When Both Variables Are Continuous

Model:  $Y = \beta_0 + \beta_1 \text{time} + \beta_2 \text{food} + \beta_3 \text{time} \times \text{food} + \epsilon$

The effect of time now depends on food amount:

$$\frac{\partial \text{weight}}{\partial \text{time}} = \beta_1 + \beta_3 \times \text{food}$$

```
# Adding hypothetical food amounts
cw <- cw %>%
  mutate(food_g = sample(c(10, 20, 30, 40, 50),
                          size = n(), replace = TRUE),
         weight2 = weight - 25 + food_g * 0.5 +
                     food_g * time * 0.1)
```

# Model 3: Multiple Interactions I



```
m_cw_3 <- gam(formula = weight2 ~  
  diet +  
  time * food_g +  
  time:diet +  
  s(chick, bs = 're'),  
  family = gaussian(),  
  data = cw,  
  method = 'ML')  
summary(m_cw_3)
```

Formula:  
weight2 ~ diet + time \* food\_g + time:diet + s(chick, bs = "re")

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	9.63434	7.15976	1.346	0.17900
diet2	-2.44531	10.26990	-0.238	0.81189
diet3	-13.19594	10.25853	-1.286	0.19889
diet4	-0.15644	10.27468	-0.015	0.98786
time	6.47388	0.42084	15.383	< 2e-16 ***
food_g	0.38868	0.14217	2.734	0.00647 **
time:food_g	0.10847	0.01168	9.290	< 2e-16 ***
diet2:time	1.86297	0.43108	4.322	1.85e-05 ***
diet3:time	4.71152	0.42929	10.975	< 2e-16 ***
diet4:time	2.94276	0.43607	6.748	3.95e-11 ***

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(chick)	41.15	46	9.854	<2e-16 ***

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.921 Deviance explained = 92.8%  
-ML = 2743.7 Scale est. = 645.31 n = 578

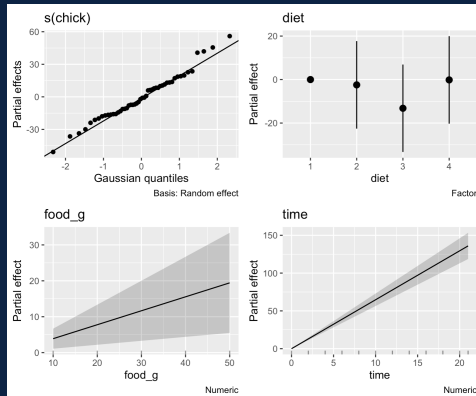
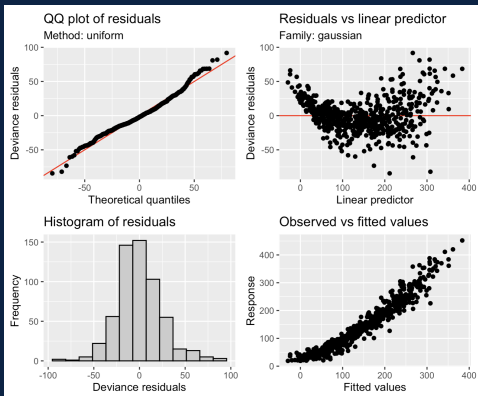
The \* operator includes main effects and interaction, while : is just the interaction

# Model 3: Multiple Interactions II



```
appraise(m_cw_3)
```

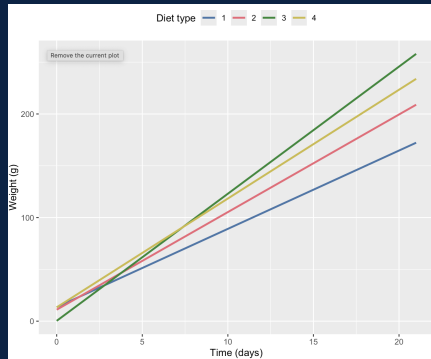
```
draw(m_cw_3, parametric = TRUE)
```



# Visualizing Continuous × Continuous Interactions I



```
preds_3 <- m_cw_3 %>%  
  data_slice(time = evenly(time),  
             diet = evenly(diet),  
             food_g = evenly(food_g),  
             chick = 'new') %>%  
  fitted_values(object = m_cw_3, data = .)  
  
# plot weight over time by diet  
preds_3 %>%  
  filter(food_g == min(food_g)) %>%  
  ggplot(aes(time, .fitted, color = diet)) +  
  geom_line(linewidth = 1) +  
  labs(x = 'Time (days)', y = 'Weight (g)') +  
  khroma::scale_color_bright(name = 'Diet type') +  
  theme(legend.position = 'top')
```




Brighter colors = higher weight; steeper gradient = stronger interaction

# Visualizing Continuous x Continuous Interactions II



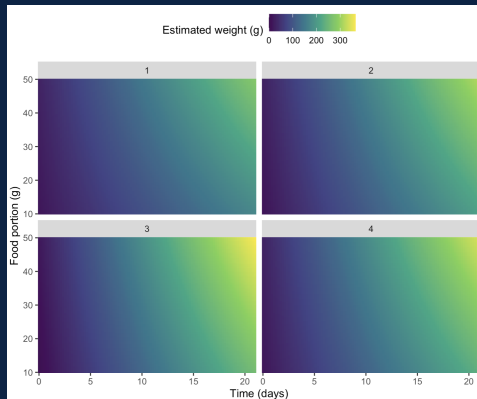
```
# plot weight over time by food_g
m_cw_3 %>%
  data_slice(time = evenly(time),
             diet = evenly(diet),
             food_g = unique(food_g),
             chick = 'new') %>%
  fitted_values(object = m_cw_3, data = .) %>%
  filter(diet == 1) %>%
  filter(food_g %in% sample(unique(food_g), 4)) %>%
  ggplot(aes(time, .fitted, group = food_g)) +
  geom_line(linewidth = 2) +
  geom_line(aes(color = food_g), linewidth = 1) +
  labs(x = 'Time (days)', y = 'Weight (g)') +
  scale_color_distiller('Food (g)', type = 'seq',
                       palette = 1,
                       direction = 1) +
  theme(legend.position = 'top')
```



# Visualizing Continuous × Continuous Interactions III



```
#' plot `time * food_g` interaction for each diet
ggplot(preds_3, aes(time, food_g, fill = .fitted)) +
  facet_wrap(~ diet) +
  geom_raster() +
  scale_fill_viridis_c('Estimated weight (g)',
                      limits = c(0, NA)) +
  scale_x_continuous('Time (days)',
                    expand = c(0, 0)) +
  scale_y_continuous('Food portion (g)',
                    expand = c(0, 0)) +
  theme(legend.position = 'top')
```



So what model does the best at predicting new observations? We haven't really examined this until now.

```
compTab <- data.frame(  
  Model = c("M0: Time only",  
            "M1: +Diet +Chick",  
            "M2: +Interactions",  
            "M3: +Continuous interactions"),  
  AIC = c(AIC(m_cw_0),  
          AIC(m_cw_1),  
          AIC(m_cw_2),  
          AIC(m_cw_3))  
)  
compTab
```

	Model	AIC
1	M0: Time only	5876.799
2	M1: +Diet +Chick	5549.704
3	M2: +Interactions	5427.771
4	M3: +Continuous interactions	5431.008

```
set.seed(10212025)
keep <- sample(1:nrow(cw), 0.8*nrow(cw), replace = F)
trainDat <- cw[keep,]; testDat <- cw[-keep, ]
m0_train <- update(m_cw_0, data = trainDat)
m1_train <- update(m_cw_1, data = trainDat)
m2_train <- update(m_cw_2, data = trainDat)
m3_train <- update(m_cw_3, data = trainDat)
testDat <- testDat %>% mutate(
  pred_m0 = predict(m0_train, newdata = testDat),
  pred_m1 = predict(m1_train, newdata = testDat),
  pred_m2 = predict(m2_train, newdata = testDat),
  pred_m3 = predict(m3_train, newdata = testDat))
rmse <- function(actual, predicted) sqrt(mean((actual - predicted)^2))
rmseSummary <- data.frame(Model = c("M0", "M1", "M2", "M3"),
  Train_R2 = c(summary(m0_train)$r.sq, summary(m1_train)$r.sq,
    summary(m2_train)$r.sq, summary(m3_train)$r.sq),
  Test_RMSE = c(rmse(testDat$weight, testDat$pred_m0),
    rmse(testDat$weight, testDat$pred_m1),
    rmse(testDat$weight, testDat$pred_m2),
    rmse(testDat$weight, testDat$pred_m3)))

rmseSummary
```

	Model	Train_R2	Test_RMSE
1	M0	0.7203430	42.21076
2	M1	0.8467806	30.65102
3	M2	0.8779115	27.76259
4	M3	0.9240923	46.15828



## Key Points for Interpretation

- **Main effects:** Average effect when other variable = 0
- **Interaction coefficient:** How much the slope changes
- **Centering helps:** Consider centering continuous variables
- **Plot, plot, plot:** Visualizations are very helpful

## Good Practice

- Include main effects with interactions
- Check if interaction improves fit
- Visualize at multiple levels

## Common Mistakes

- Interaction without main effects
- Over-interpreting small interactions
- Ignoring multicollinearity

## Exercise: Explore the sleep study data

```
# load the sleep study data which is a part of the package 'lme4'
library('lme4')
?sleepstudy  # reaction times with sleep deprivation

# Your tasks:
# 1. Plot reaction time vs days for each subject
# 2. Fit a model with random intercepts for subjects
# 3. Add random slopes (do subjects differ in how sleep loss affects them?)
# 4. Compare models and determine a 'best' model using diagnostics

# you might wish to start with
ggplot(sleepstudy, aes(Days, Reaction)) +
  geom_line(aes(group = Subject)) +
  facet_wrap(~ Subject)
```

- Interactions allow effects to vary by context:  $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$
- Fixed effects: specific levels we care about
- Random effects: sample from population, allow pooling
- Hierarchical models handle grouped/nested data
- Usually you want to include main effects with interactions
- Visualization is important for understanding interactions

**Remember** that with interactions, "it depends" becomes part of your model!

Additional Questions?  
Book an Appointment!



## Next Workshop:

Generalized Linear Models October

23, 4:00 PM

- Beyond normal distributions
- Link functions

## Additional Resources:

- *Data Analysis Using Regression and Multilevel/Hierarchical Models* (Gelman & Hill)