



# Penalized Models

Fitting Models to Data, Not Data to Models

Model Fitting Series - With Applications in R

---

Jesse Ghashti

Source code by Stefano Mezzini

November 18, 2025

Centre for Scholarly Communication

The University of British Columbia | Okanagan Campus | Syilx Okanagan Nation Territory

# Workshop Series Overview



Session	Topic	Date/Time
1	Simple Linear Regression	Oct 7, 9:00 AM
2	Fitting Linear Models in R	Oct 8, 10:30 AM
3	Multiple Linear Regression in R	Oct 16, 4:00 PM
4	Interaction Terms & Hierarchical Linear Models	Oct 21, 11:00 AM
5	Generalized Linear Models	Oct 23, 4:00 PM
6	Generalized Additive Models (GAMs)	Oct 28, 11:00 AM
7	Interpreting & Predicting from GAMs	Oct 29, 10:30 AM
8	Hierarchical GAMs	Nov 4, 12:00 PM
9	Penalized Models	Nov 18, 11:00 AM
10	Survival Models	Nov 25, 11:00 AM
11	Nonparametric Models	Dec 2, 11:00 AM



**New to R?** Check out the Fundamentals of R series!

**GitHub code and slides** for today's workshop (and previous workshops)



Alternatively, code/slides available at the bottom of  
<https://csc-ubc-okanagan.github.io/workshops/>



## Key Concepts

- Reviewed smooth types in `mgcv`: `tp`, `cr`, `cc`, `ad`
- Learned factor smooths (`fs`, `sz`) and random effects (`re`, `mrf`)
- Built interactions with `s()`, `te()`, and `ti()`
- Applied to seasonal  $\times$  long-term trends in temperature data
- Key insight: `ti()` separates main effects from interactions!

Today we will...

- Understand why we need regularization (the bias-variance tradeoff)
- Learn Ridge, Lasso, and Elastic Net penalties
- Fit penalized models with `glmnet`
- Choose tuning parameters via cross-validation
- Compare variable selection approaches
- Apply to high-dimensional data (many predictors)

Today we require...

```
library('dplyr')      # data wrangling
library('tidyverse')   # data reshaping
library('ggplot2')     # plotting
library('glmnet')      # penalized regression
library('ISLR')        # example datasets
theme_set(theme_bw(base_size = 15))
```



## When Standard Regression Fails

### Issues with many predictors:

- **Overfitting:** Model fits training data perfectly but predicts poorly
- **Multicollinearity:** Correlated predictors  $\Rightarrow$  unstable coefficients
- **High variance:** Small changes in data  $\Rightarrow$  large coefficient changes
- **$p > n$  problem:** More predictors than observations  $\Rightarrow$  no unique solution

### What we want:

- Stable coefficient estimates
- Better out-of-sample predictions
- Variable selection (which predictors matter?)
- Interpretable models



## Prediction Error Decomposition

### Expected prediction error:

$$\text{MSE} = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}$$

### Unpenalized regression:

- Low bias (fits training data well)
- High variance (sensitive to data changes)
- Poor generalization to new data

### Penalized regression:

- Slightly higher bias (shrinks coefficients)
- Much lower variance (more stable)
- Better overall prediction error

## Objective Function

**Standard least squares:**

$$\hat{\beta} = \arg \min_{\beta} \sum_{i=1}^n (y_i - x_i^T \beta)^2$$

**Penalized least squares:**

$$\hat{\beta}_{\lambda} = \arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - x_i^T \beta)^2 + \lambda \cdot \text{Penalty}(\beta) \right\}$$

where:

- $\lambda \geq 0$  is the **tuning parameter** (controls strength)
- Larger  $\lambda \Rightarrow$  more shrinkage toward zero
- $\lambda = 0 \Rightarrow$  ordinary least squares

## Ridge Penalty

### Objective:

$$\hat{\beta}^{\text{ridge}} = \arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - x_i^T \beta)^2 + \lambda \sum_{j=1}^p \beta_j^2 \right\}$$

### Properties:

- Penalizes the sum of *squared* coefficients
- Shrinks all coefficients toward zero
- **Does NOT set coefficients exactly to zero**
- Good for multicollinearity
- Keeps all predictors in the model

You might consider using ridge regression when all predictors might be relevant, or you have severe multicollinearity

## Lasso Penalty

### Objective:

$$\hat{\beta}^{\text{lasso}} = \arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - x_i^T \beta)^2 + \lambda \sum_{j=1}^p |\beta_j| \right\}$$

### Properties:

- Penalizes the sum of *absolute values* of coefficients
- Shrinks coefficients toward zero
- **Can set coefficients exactly to zero** (variable selection!)
- Produces sparse models (fewer predictors)
- Performs automatic feature selection

You might consider using lasso regression when you want a simpler model and believe many predictors are irrelevant



## Elastic Net Penalty

### Objective:

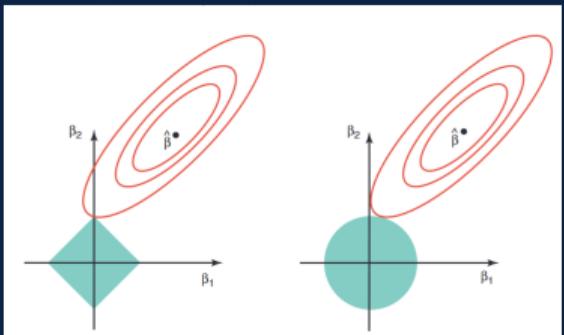
$$\hat{\beta}^{\text{enet}} = \arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - x_i^T \beta)^2 + \lambda \left( \alpha \sum_{j=1}^p |\beta_j| + (1 - \alpha) \sum_{j=1}^p \beta_j^2 \right) \right\}$$

### Properties:

- Combines Ridge and Lasso penalties
- $\alpha \in [0, 1]$  controls the mix:  $\alpha = 1$  (Lasso),  $\alpha = 0$  (Ridge)
- Can select groups of correlated variables together
- More stable than Lasso when  $p > n$

You might consider using Elastic Net when you want variable selection but have correlated predictors

# Geometric Intuition: Why Lasso Gives Sparse Solutions



**Ridge:** Circular constraint  $\Rightarrow$  solution rarely at axis

**Lasso:** Diamond constraint  $\Rightarrow$  corners at axes  $\Rightarrow$  exact zeros

# Exploring the Diabetes Dataset

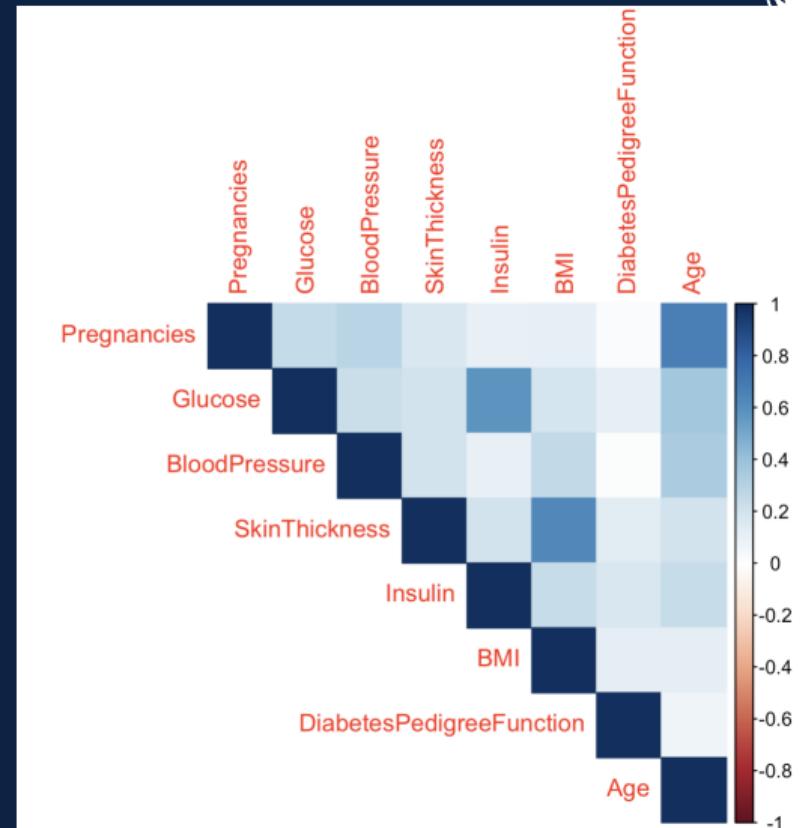


```
library(ISLR)
library(corrplot)

data(diabetes)
Diabetes <- na.omit(diabetes) # remove missing
dim(Diabetes)

# Check correlations
numVars <- Diabetes[, sapply(Diabetes, is.numeric)]
numVars <- numVars[, !colnames(numVars) %in% "Outcome"]
corMat <- cor(numVars) # get correlations
corrplot(corMat, method = 'color',
         type = 'upper') # plot correlations
```

Many predictors are highly correlated  
— multicollinearity...



# Preparing Data for glmnet



```
# glmnet requires MATRIX input (not data frame)
# model.matrix automatically creates dummy variables for factors
X <- model.matrix(Outcome ~ ., data = Diabetes)[, -1] # remove intercept
y <- Diabetes$Outcome

# split into 70/30 train/test set
set.seed(2025)
trainIDX <- sample(1:nrow(X), size = floor(0.7 * nrow(X)))
testIDX <- setdiff(1:nrow(X), trainIDX)
Xtrain <- X[trainIDX, ]
ytrain <- y[trainIDX]
Xtest <- X[testIDX, ]
ytest <- y[testIDX]
```

glmnet requires matrix input and automatically standardizes predictors.

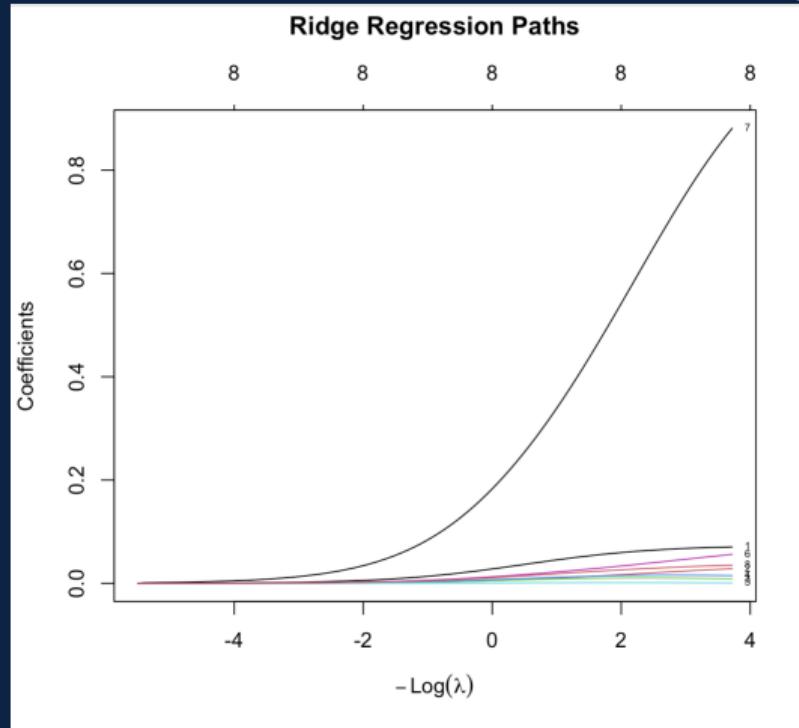
# Fitting Ridge Regression



```
# Fit ridge regression
# alpha = 0 for ridge

ridgeModel <- glmnet(X, y, alpha = 0, family = "binomial")
plot(ridgeModel, xvar = 'lambda', label = TRUE)
title('Ridge Regression Paths', line = 3)

# coefficients at specific lambdas
coef(ridgeModel, s = 50)
coef(ridgeModel, s = 500)
coef(ridgeModel, s = 5000)
# Take a look at these :)
```

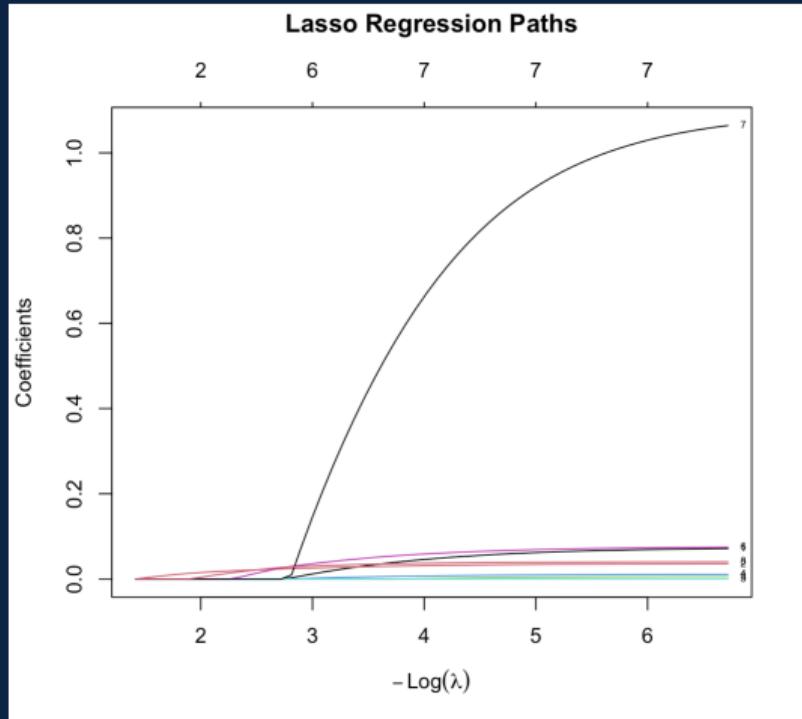


# Fitting Lasso Regression



```
# Fit lasso regression
# alpha = 1 for lasso
lassoModel <- glmnet(X, y, alpha = 1, family = "binomial")
plot(lassoModel, xvar = 'lambda', label = TRUE)
title('Lasso Regression Paths', line = 3)

# coefficients at specific lambdas (play around with these!)
coef(lassoModel, s = 50)
coef(lassoModel, s = 500)
coef(lassoModel, s = 5000)
# Notice that many coefficients are
# EXACTLY zero (which no longer contribute to the model)
```



# Cross-Validation: Choosing $\lambda$ for Ridge



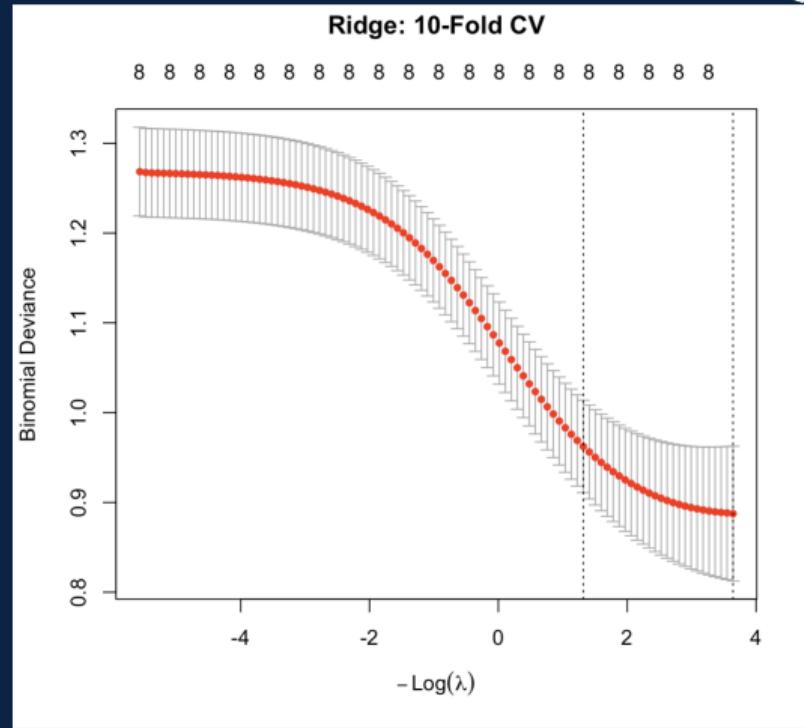
```
# 10-fold cross-validation for Ridge
set.seed(2025)
cvRidge <- cv.glmnet(Xtrain,
                      ytrain,
                      alpha = 0,
                      nfolds = 10,
                      family = "binomial")

plot(cvRidge)
title('Ridge: 10-Fold CV', line = 3)

# best lambda values
cvRidge$lambda.min # min CV error is 0.02615955
cvRidge$lambda.1se # 1 SE rule is 0.2677511

# optimal coefficients
coef(cvRidge, s = 'lambda.min')
```

**lambda.min:** Minimizes CV error  
**lambda.1se:** Simpler, more stable



# Cross-Validation: Choosing $\lambda$ for Lasso



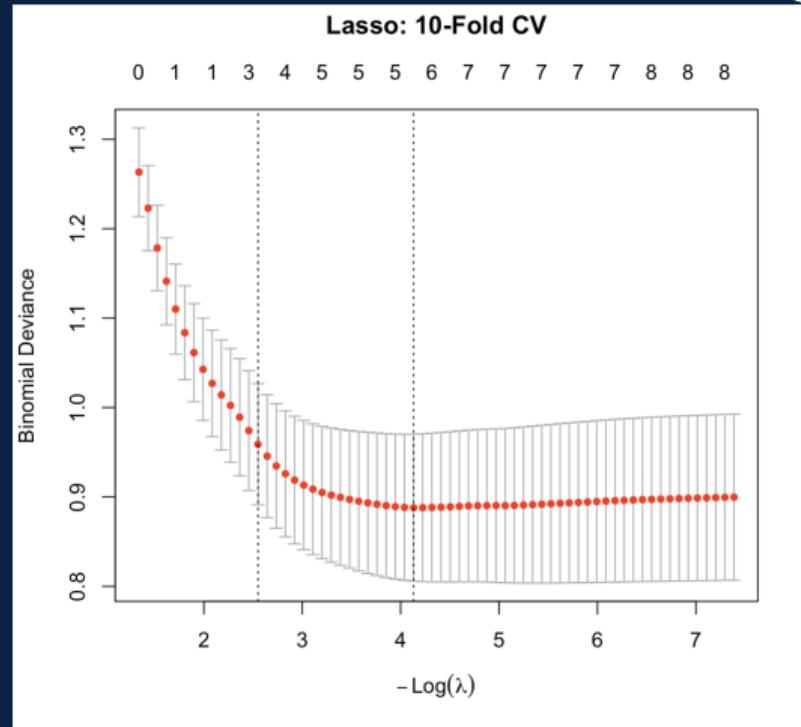
```
# 10-fold cross-validation for Lasso
set.seed(2025)
cvLasso <- cv.glmnet(Xtrain,
                      ytrain,
                      alpha = 1,
                      nfolds = 10,
                      family = "binomial")

plot(cvLasso)
title('Lasso: 10-Fold CV', line = 3)

# best lambda values
cvLasso$lambda.min # 0.01605126
cvLasso$lambda.1se # 0.07805087

# optimal coefficients
coef(cvLasso, s = 'lambda.min')
```

Number at top shows how many variables included



# Making Predictions and Comparing Models



```
# predict on test set using optimal lambda
ridgePredict <- predict(cvRidge, s = 'lambda.min', newx = Xtest, type = "response")
lassoPredict <- predict(cvLasso, s = 'lambda.min', newx = Xtest, type = "response")

# we will compare to multiple linear regression as well.
lmModel <- glm(Salary ~ ., data = Hitters[trainIDX, ], family = "binomial")
lmPredict <- predict(lmModel, newdata = Hitters[testIDX, ], type = "response")]

accuracy <- function(actual, predicted) {
  predicted <- ifelse(predicted >= 0.5, 1, 0)
  matches <- sum(diag(Threshersh::matchLabels(table(actual, predicted))))
  return(matches/length(predicted))
}

accuracyLM <- accuracy(as.numeric(ytest), lmPredict)
accuracyRIDGE <- accuracy(as.numeric(ytest), ridgePredict)
accuracyLASSO <- accuracy(as.numeric(ytest), lassoPredict)

data.frame(
  Model = c('LM', 'Ridge', 'Lasso'),
  accuracy = c(accuracyLM, accuracyRIDGE, accuracyLASSO)
)
```

# Test Set Performance Comparison



Model	Test Accuracy	# Variables
LM	0.7223	9
Ridge	0.7426	9
Lasso	0.7223	6

Lasso does better as good as LM with predictions but with fewer variables (simpler, more interpretable)

Both penalized methods outperform LM on out-of-sample data, but for different reasons

# Variable Selection: Which Predictors Matter?



```
lassoCoefficients <- coef(cvLasso,
                           s = 'lambda.min')
lassoCoefficients
# which variables selected?
varsSelected <-
  lassoCoefficients[lassoCoefficients[, 1] != 0, ]
varsSelected
```

(Intercept)	Pregnancies
-9.80932850	0.05202342
Glucose	BMI
0.03935407	0.08221258
DiabetesPedigreeFunction	Age
0.54386344	0.02382367

# Elastic Net: Combining Ridge and Lasso



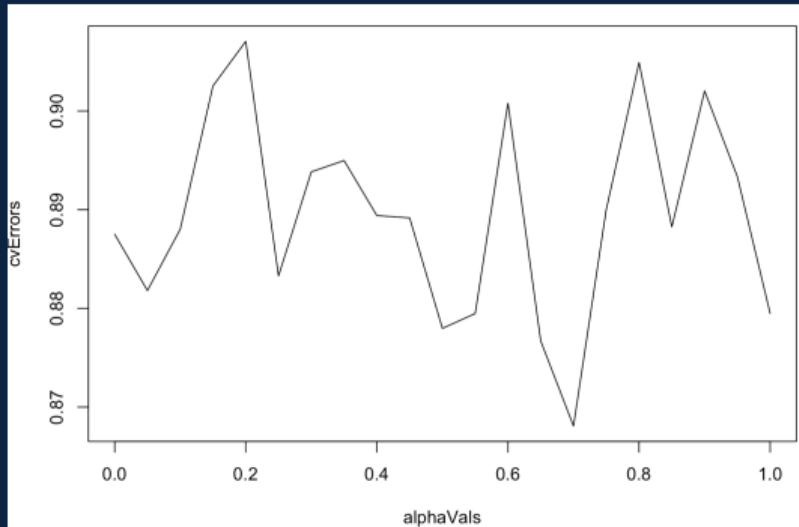
```
set.seed(2025)
cvENET <- cv.glmnet(Xtrain, ytrain, alpha = 0.5, nfolds = 10, family = "binomial")
plot(cvENET)
title('Elastic Net ( = 0.5): 10-Fold CV', line = 3)
enetPredict <- predict(cvENET, s = 'lambda.min', newx = Xtest, type = "response")
accuracyENET <- accuracy(as.numeric(ytest), enetPredict)
data.frame(
  Model = c('LM', 'Ridge', 'Lasso', 'Elastic Net'),
  accuracy = c(accuracyLM, accuracyRIDGE, accuracyLASSO, accuracyENET)
)

# how many variables selected?
enetCoefficients <- coef(cvENET, s = 'lambda.min')
sum(enetCoefficients[-1] != 0) # 7 selected
```

# Tuning $\alpha$ : Finding the Best Mix



```
alphaVals <- seq(0, 1, by = 0.05)
cvResults <- list()
set.seed(2025)
for (i in seq_along(alphaVals)) {
  cvResults[[i]] <-
    cv.glmnet(Xtrain, ytrain,
               alpha = alphaVals[i],
               nfolds = 10,
               family = "binomial")
}
# minimum CV error
cvErrors <- sapply(cvResults,
                     function(x) min(x$cvm))
alphaOptimal <-
  alphaVals[which.min(cvErrors)]
alphaOptimal # 0.7
```



# Final Model Comparison



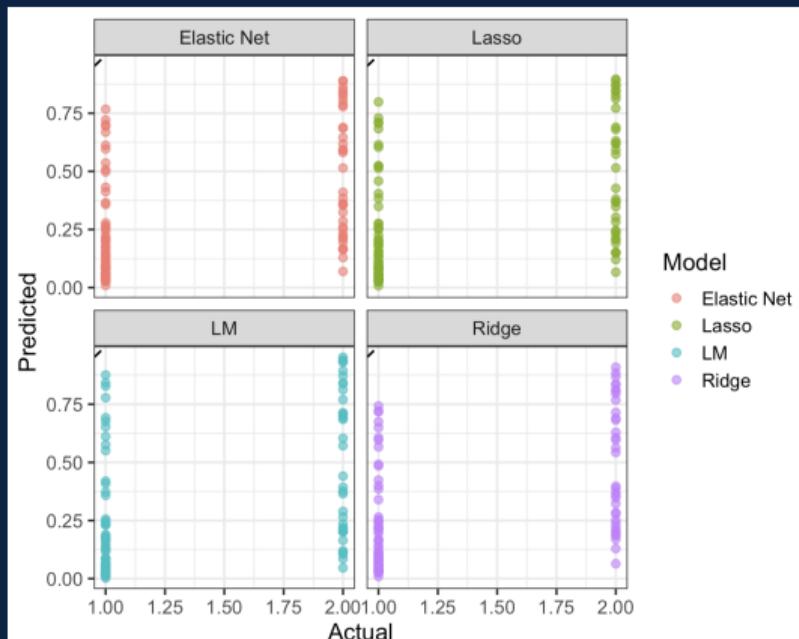
Model	Test Accuracy	# Variables	$\alpha$
LM	0.7228	9	N/A
Ridge	0.7426	9	0
Lasso	0.7228	6	1
Elastic Net ( $\alpha=0.5$ )	0.7327	7	0.5
Optimal ( $\alpha=0.7$ )	0.7228	5	0.7

Optimal  $\alpha \approx 0.7$ : Closer to Lasso, emphasizing variable selection

# Visualizing Predictions: All Models



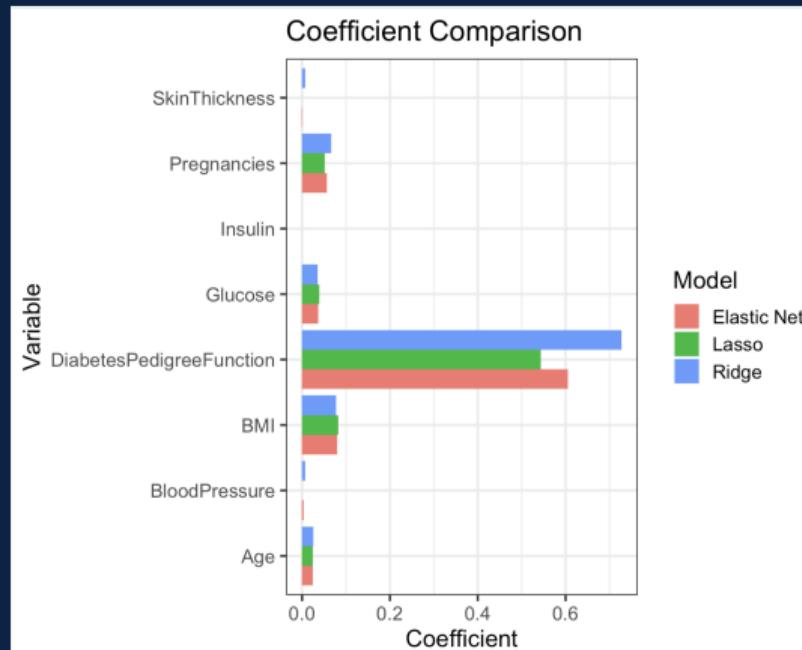
```
predictionCompare <- tibble(  
  Actual = as.numeric(ytest),  
  LM = as.vector(lmPredict),  
  Ridge = as.vector(ridgePredict),  
  Lasso = as.vector(lassoPredict),  
  `Elastic Net` = as.vector(enetPredict)  
) %>%  
  pivot_longer(cols = -Actual,  
              names_to = 'Model',  
              values_to = 'Predicted')  
  
ggplot(predictionCompare,  
       aes(Actual, Predicted,  
            color = Model)) +  
  geom_point(alpha = 0.6) +  
  geom_abline(slope = 1,  
              intercept = 0,  
              linetype = 'dashed') +  
  facet_wrap(~ Model)
```



# Coefficient Comparison Across Methods



```
coefficientComparison <- tibble(
  Variable = rownames(lassoCoefficients)[-1],
  Ridge = as.vector(coef(cvRidge, s = 'lambda.min')[-1]),
  Lasso = as.vector(lassoCoefficients[-1]),
  `Elastic Net` = as.vector(enetCoefficients[-1])
) %>%
  pivot_longer(cols = -Variable,
               names_to = 'Model',
               values_to = 'Coefficient')
ggplot(coefficientComparison,
       aes(Variable, Coefficient,
            fill = Model)) +
  geom_col(position = 'dodge') +
  coord_flip() +
  labs(title = 'Coefficient Comparison')
```



# High-Dimensional Data: $p > n$



```
# simulate high-dimensional data where p > n
set.seed(2025)
n <- 100      # observations
p <- 200      # predictors

# only first 10 predictors are truly important
Xhd <- matrix(rnorm(n * p), n, p)
trueCoefficients <- c(rep(2, 5), rep(-2, 5), rep(0, p - 10))
yhd <- Xhd %*% trueCoefficients

# try the linear model, it should fail!
# lmHD <- lm(yhd ~ Xhd)

set.seed(2025)
cvHD <- cv.glmnet(Xhd, yhd, alpha = 1, nfolds = 10)

# how many variables selected?
coef_hd <- coef(cvHD, s = 'lambda.1se')
sum(coef_hd[-1] != 0)

# which variables were selected? (hopefully the first 10)
which(coef_hd[-1] != 0)
```

# High-Dimensional Lasso ( $p > n$ )



```
> sum(coef_hd[-1] != 0)
[1] 11
> # which variables were selected? (hopefully the first 10)
> which(coef_hd[-1] != 0)
[1]  1  2  3  4  5  6  7  8  9 10 165
```

Lasso identifies relevant variables even when  $p > n$  (impossible for LM)



## The One Standard Error Rule

### `lambda.min`:

- Minimizes cross-validation error
- Best predictive performance
- May overfit slightly
- More variables selected (less sparse)

### `lambda.1se`:

- Largest  $\lambda$  within 1 SE of minimum
- More regularization (simpler model)
- Better generalization in practice
- Fewer variables (more sparse)
- Preferred for interpretability

Rule of thumb: Use `lambda.1se` for simpler, more stable models

## Important Details

### Standardization:

- `glmnet` standardizes by default (`standardize = TRUE`)
- Essential for fair penalization across different scales
- Final coefficients returned on original scale

### Categorical predictors:

- Use `model.matrix()` to create dummy variables
- `glmnet` doesn't handle factors directly
- Can group dummies with `penalty.factor`

### Response transformations:

- For GLM families: set `family` argument
- Logistic: `family = 'binomial'`
- Poisson: `family = 'poisson'`

# Decisions, Decisions, Decisions



Consideration	Recommendation
Severe multicollinearity	Ridge (keeps all variables)
Want interpretable model	Lasso (variable selection)
$p > n$ problem	Lasso or Elastic Net
Correlated predictors + selection	Elastic Net ( $\alpha \approx 0.5 - 0.7$ )
All predictors relevant	Ridge
Many irrelevant predictors	Lasso
Uncertain	Tune $\alpha$ via CV

## Do's

- Always use cross-validation to choose  $\lambda$
- Consider `lambda.1se` for simpler, more stable models
- Standardize predictors (`glmnet` does this automatically)
- Split data into train/test for honest evaluation
- Try multiple  $\alpha$  values for Elastic Net
- Use `model.matrix()` to handle categorical predictors

## Don'ts

- Don't use a single  $\lambda$  without CV
- Don't forget to check test set performance
- Don't interpret Ridge coefficients as "importance"
- Don't ignore the geometric interpretation



## Unified Framework

### Remember from Workshop 6-8:

- GAMs use penalization on smooth functions
- Penalty controls "wiggliness" of curves
- REML/ML chooses optimal smoothing parameter

### Penalized regression:

- Similar idea applied to coefficients
- Penalty controls "size" of coefficients
- Cross-validation chooses optimal  $\lambda$

### Common theme:

- Bias-variance tradeoff
- Automatic complexity control
- Better out-of-sample prediction

# You Try: Penalized Regression Practice



## Exercise: Compare methods

```
# Use the College dataset from ISLR
library(ISLR)
data(College)
?College

# Task: Predict Graduation Rate (Grad.Rate) from all other variables

# 1. Prepare data
#     - Create X matrix and y vector
#     - Split into train (70%) and test (30%)

# 2. Fit three models
#     - Ridge (alpha = 0)
#     - Lasso (alpha = 1)
#     - Elastic Net (alpha = 0.5)
#     - Use CV to find optimal lambda for each

# 3. Compare performance
#     - Calculate test RMSE for each
#     - How many variables does Lasso select?
#     - Which model performs best?

# 4. Interpret
#     - Plot coefficient paths
#     - Which predictors are most important?
#     - Try different alpha values - does it matter?
```

# Key Takeaways



- Penalized regression addresses overfitting, multicollinearity, and  $p > n$
- Ridge shrinks coefficients but keeps all variables
- Lasso performs variable selection (sets coefficients to zero)
- Elastic Net combines both approaches (tune  $\alpha$ )
- Always use cross-validation to choose  $\lambda$
- Consider `lambda.1se` for simpler, more stable models
- Penalized methods usually outperform LM on test data
- The bias-variance tradeoff is central to statistical learning
- `glmnet` is fast, efficient, and works when  $p > n$
- Geometric intuition explains why Lasso gives sparse solutions



Additional Questions?  
Book an Appointment!



Next Workshop:

**Survival Models**

→ November 25, 11:00 AM

# Thank You!

## Questions?

**Workshop Materials:**

<https://github.com/csc-ubc-okanagan/ubco-csc-modeling-workshop>

**Contact:**

[jesse.ghashti@ubc.ca](mailto:jesse.ghashti@ubc.ca)