



Survival Models

Fitting Models to Data, Not Data to Models

Model Fitting Series - With Applications in R

Jesse Ghashti

Source code by Jesse Ghashti

November 25, 2025

Centre for Scholarly Communication

The University of British Columbia | Okanagan Campus | Syilx Okanagan Nation Territory

Workshop Series Overview



Session	Topic	Date/Time
1	Simple Linear Regression	Oct 7, 9:00 AM
2	Fitting Linear Models in R	Oct 8, 10:30 AM
3	Multiple Linear Regression in R	Oct 16, 4:00 PM
4	Interaction Terms & Hierarchical Linear Models	Oct 21, 11:00 AM
5	Generalized Linear Models	Oct 23, 4:00 PM
6	Generalized Additive Models (GAMs)	Oct 28, 11:00 AM
7	Interpreting & Predicting from GAMs	Oct 29, 10:30 AM
8	Hierarchical GAMs	Nov 4, 12:00 PM
9	Penalized Models	Nov 18, 11:00 AM
10	Survival Models	Nov 25, 11:00 AM
11	Nonparametric Models	Dec 2, 11:00 AM



New to R? Check out the Fundamentals of R series!

GitHub code and slides for today's workshop (and previous workshops)



Alternatively, code/slides available at the bottom of
<https://csc-ubc-okanagan.github.io/workshops/>



Today we will...

- Build a realistic **simulated reliability dataset** (pumps in plants)
- Fit and interpret **Cox PH models** with and without plant-level frailty
- Allow for **time-varying covariate effects** (non-proportional hazards)
- Compare **parametric AFT models** (Weibull, lognormal) with Cox
- Fit a **piecewise exponential model** via Poisson regression
- Evaluate **predictive accuracy**: AIC, C-index, prediction error curves

We will use these packages...

```
# install.packages(c("survival", "survminer", "dplyr", "ggplot2", "broom", "pec"))
library(survival)
library(survminer)
library(dplyr)
library(ggplot2)
library(broom)
library(pec)
```

Reliability Setup

We simulate time-to-failure for industrial pumps installed in **3 plants**:

- Each pump belongs to a plant $\in \{P1, P2, P3\}$
- **Continuous covariates:**
 - Temperature (temp): degrees C above baseline
 - Load (load): relative design load
- **Categorical covariates:**
 - Material (material = A/B)
 - Maintenance program (maint = 0/1)
- **Right censoring** between 40 and 120 months
- True model has a **plant-level frailty** (unobserved plant stressors)

Simulating Pump Reliability Data



Simulation code

```
## Simulated Data
## Features:
## 1. 3 plants with different baseline reliability (cluster effect)
## 2. Continuous covariates: temperature, load
## 3. Categorical: material (A/B), maintenance program (0/1)
## Right censoring, true model has a plant-level random effect (frailty)

numberOFplants <- 3
plant      <- factor(sample(paste0("P", 1:numberOFplants), size = 500, replace = TRUE))
temp       <- rnorm(500, mean = 25, sd = 4)           # degrees C above baseline
load       <- runif(500, 0.7, 1.3)                    # design load
material   <- factor(sample(c("A", "B"), 500, TRUE)) # material it was made from
maint     <- rbinom(500, 1, 0.5)                     # maintenance program (think of regular or sparse)

## plant frailty: log-normal random effect on hazard. this can be thought of as unobserved plant-level stressors
plantsEFFECT <- rnorm(numberOFplants, mean = 0, sd = 0.5)
names(plantsEFFECT) <- paste0("P", 1:numberOFplants)
## baseline hazards and their respective coefficients
lambda0 <- 0.005          # baseline rate
baselineTEMP <- 0.04        # per °C
baselineLOAD <- 1.2
baselineMATERIALB <- -0.5
baselineMAINTENANCE <- -0.7
```

Simulating Pump Reliability Data



Simulation code

```
## predictor for the hazard occurring
linpred <- baselineTEMP*(temp-mean(temp)) +
  baselineLOAD*(load-1) +
  baselineMATERIALB*(material=="B") +
  baselineMAINTENANCE*maint +
  plantsEFFECT[plant]

lambdaI <- lambda0*exp(linpred)

## we will model with exponential survival times
U <- runif(500)
trueTIME <- -log(U)/lambdaI
## Random censoring between 40 and 120 months
censorTIME <- runif(500, 40, 120)
time <- pmin(trueTIME, censorTIME)
status <- as.integer(trueTIME <= censorTIME) # 1=failure, 0=censored

dat <- data.frame(pump_id = 1:500, plant,
  time, status, temp, load, material, maint =
  factor(maint))
head(dat)
```

	pump_id	plant	time	status	temp	load	material	maint
1	1	P1	38.34740	1	28.34623	1.2570109	A	0
2	2	P1	35.70169	1	26.16657	1.2670218	A	1
3	3	P1	86.44448	0	24.45903	0.9121598	B	0
4	4	P1	60.23994	1	22.21728	1.1337975	A	0
5	5	P1	40.82726	0	24.54467	0.9440624	B	1

Nonparametric survival

We first explore survival differences **across plants**:

- Construct $S(t)$ using **Kaplan–Meier** for each plant
- Check:
 - Which plant has higher survival (better reliability)?
 - Are differences mostly early or late in follow-up?
- This is a model-free, descriptive step before regression


KM Curves by Plant in R



Kaplan-Meier

```
## KM by plant
survivalOBJECT <- Surv(dat$time, dat$status)
kmPLANT <- survfit(survivalOBJECT ~ plant,
                     data = dat)

ggsurvplot(
  kmPLANT,
  conf.int    = TRUE,
  risk.table = TRUE,
  ggtheme     = theme_minimal(),
  xlab        = "Time since installation (months)",
  ylab        = "Survival probability",
  title       = "Kaplan-Meier Curves by Plant"
)
```



Cox PH

Model...

$$\lambda(t \mid X) = \lambda_0(t) \exp(\beta^\top X)$$

- $\lambda_0(t)$: unspecified baseline hazard
- β : log hazard ratios for covariates
- Proportional hazards assumption: hazard ratios do not change over time

To interpret...

- $\exp(\beta_j)$ is the hazard ratio for a 1-unit increase in covariate X_j
- > 1 : higher failure rate (shorter survival); < 1 : lower failure rate (longer survival)

Cox PH Model for Pump Data



Fit Cox model

```
## Cox PH model
coxBASIC <- coxph(
  Surv(time, status) ~ scale(temp) +
    scale(load) +
    material+maint,
  data = dat,
  x = TRUE,
  y = TRUE
)
summary(coxBASIC)
broom::tidy(coxBASIC, exponentiate = TRUE,
            conf.int = TRUE)
```

```
Call:
coxph(formula = Surv(time, status) ~ scale(temp) + scale(load) +
material + maint, data = dat, x = TRUE, y = TRUE)

n= 500, number of events= 89

      coef exp(coef) se(coef)     z Pr(>|z|)
scale(temp) -0.05947   0.94226  0.10591 -0.562  0.57442
scale(load)  0.32119   1.37877  0.11196  2.869  0.00412 **
materialB   -0.38906   0.67769  0.21403 -1.818  0.06910 .
maint1     -0.71371   0.48983  0.22512 -3.170  0.00152 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

      exp(coef) exp(-coef) lower .95 upper .95
scale(temp)   0.9423    1.0613   0.7656   1.1596
scale(load)   1.3788    0.7253   1.1071   1.7171
materialB    0.6777    1.4756   0.4455   1.0309
maint1       0.4898    2.0415   0.3151   0.7615

Concordance= 0.638  (se = 0.03 )
Likelihood ratio test= 22.47 on 4 df,  p=2e-04
Wald test        = 21.13 on 4 df,  p=3e-04
Score (logrank) test = 21.8 on 4 df,  p=2e-04


> broom::tidy(coxBASIC, exponentiate = TRUE, conf.int = TRUE)
# A tibble: 4 × 7
  term      estimate std.error statistic p.value conf.low conf.high
  <chr>      <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>
1 scale(temp)  0.942     0.106    -0.562    0.574     0.766     1.16
2 scale(load)   1.38      0.112     2.87    0.00412    1.11      1.72
3 materialB    0.678      0.214    -1.82    0.0691     0.446     1.03
4 maint1       0.490      0.225    -3.17    0.00152    0.315     0.761
```

Cox PH Model for Pump Data



Check PH

```
## check the proportional hazards assumptions  
coxBASIC_zph <- cox.zph(coxBASIC)  
coxBASIC_zph  
plot(coxBASIC_zph)
```



Unobserved heterogeneity

Pumps within the same plant may share unmeasured risk factors (frailty):

$$\lambda_{ij}(t) = \lambda_0(t) \exp(\beta^\top X_{ij} + u_{k(i)})$$

- i : pump, $k(i)$: plant
- u_k : plant-level random effect (**frailty**)
- Captures unobserved plant stressors (e.g., operating conditions, management)

Why use frailty models?

- Account for within-plant clustering
- Get more realistic uncertainty and plant-specific hazard multipliers

Cox Model with Plant Frailty



Fitting frailty model

```
## Cox model with plant-level frailty
coxFRAILTY <- coxph(
  Surv(time, status) ~ scale(temp)+scale(load)
    +material+maint + frailty(plant),
  data = dat,
  x = TRUE,
  y = TRUE
)
summary(coxFRAILTY)
## compare log-likelihood/AIC with and without frailty
AIC(coxBASIC, coxFRAILTY)
```

```
coxph(formula = Surv(time, status) ~ scale(temp) + scale(load) +
  material + maint + frailty(plant), data = dat, x = TRUE,
  y = TRUE)

n= 500, number of events= 89
```

	coef	se(coef)	se2	Chisq	DF	p	
scale(temp)	-0.06023	0.1051	0.1051	0.33	1.00	0.5700	
scale(load)	0.32854	0.1120	0.1120	8.61	1.00	0.0033	
materialB	-0.39137	0.2141	0.2141	3.34	1.00	0.0680	
maint1	-0.74047	0.2257	0.2255	10.77	1.00	0.0010	
frailty(plant)					5.72	1.31	0.0260

	exp(coef)	exp(-coef)	lower	.95	upper	.95
scale(temp)	0.9415	1.0621	0.7663		1.1569	
scale(load)	1.3889	0.7200	1.1152		1.7298	
materialB	0.6761	1.4790	0.4444		1.0288	
maint1	0.4769	2.0969	0.3064		0.7422	
gamma:P1	1.2737	0.7851	0.9010		1.8005	
gamma:P2	0.7836	1.2762	0.5261		1.1671	
gamma:P3	0.9427	1.0608	0.6443		1.3793	

Iterations: 10 outer, 29 Newton-Raphson
Variance of random effect= 0.06355876 I-likelihood = -511.7
Degrees of freedom for terms= 1.0 1.0 1.0 1.0 1.3
Concordance= 0.658 (se = 0.03)
Likelihood ratio test= 30.06 on 5.3 df, p=2e-05

```
> ## compare log-likelihood/AIC with and without frailty
> AIC(coxBASIC, coxFRAILTY)
df      AIC
coxBASIC 4.000000 1033.967
coxFRAILTY 5.304768 1028.984
```


Cox Model with Plant Frailty



Frailty model coefficient visualization

```
## random effect estimates (frailties)
frailtyVALS <- data.frame(
  plant = c("p1", "p2", "p3"),
  frailty = coxFRAILTY$coefficients[c("gamma:P1",
    "gamma:P2", "gamma:P3")]
)

ggplot(frailtyVALS, aes(x = plant, y = frailty)) +
  geom_col() +
  theme_minimal() +
  labs(
    x = "Plant",
    y = "Estimated frailty (log hazard multiplier)",
    title = "Estimated Plant-Level Frailties (Cox Model)"
)
```



When PH fails

If diagnostics suggest violation of PH for a covariate (e.g. load):

- Effect of that covariate may change over time
- A simple approach: allow a time-varying coefficient

$$\lambda(t \mid X) = \lambda_0(t) \exp\{\beta^\top X + \gamma \cdot \text{load} \cdot f(t)\}$$

- Here we choose $f(t) = \log(t + 1)$
- Interpretation: the importance of load increases/decreases over log(time)

Time-Varying Effect of Load



Cox with tt()

```
## non-proportional hazards: time-varying effect
## suppose diagnostic suggested load has
## time-varying effect. We will now allow
## an effect of load that changes ~ log(time).
## tt() gets covariate*function of time;
coxTV <- coxph(
  Surv(time, status) ~ scale(temp)+material
    +maint+tt(load),
  data = dat,
  tt = function(x, t, ...) x*log(t+1)
)
summary(coxTV)
## the coefficient of tt(load) describes
## how effect of load increases/decreases
## over log(time). In reality, this is
## interpreted as load matters more/less
## as pumps age.
```

```
Call:
coxph(formula = Surv(time, status) ~ scale(temp) + material +
  maint + tt(load), data = dat, tt = function(x, t, ...) x *
  log(t + 1))

n= 500, number of events= 89

            coef exp(coef) se(coef)      z Pr(>|z|)
scale(temp) -0.05822   0.94344  0.10589 -0.550  0.58244
materialB   -0.39300   0.67503  0.21416 -1.835  0.06649 .
maint1     -0.71585   0.48878  0.22517 -3.179  0.00148 **
tt(load)     0.52372   1.68829  0.18177  2.881  0.00396 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

            exp(coef) exp(-coef) lower .95 upper .95
scale(temp)    0.9434     1.0600    0.7666    1.1610
materialB     0.6750     1.4814    0.4436    1.0271
maint1        0.4888     2.0459    0.3144    0.7599
tt(load)       1.6883     0.5923    1.1823    2.4109

Concordance= 0.639  (se = 0.031 )
Likelihood ratio test= 22.56 on 4 df,  p=2e-04
Wald test          = 21.15 on 4 df,  p=3e-04
Score (logrank) test = 21.84 on 4 df,  p=2e-04
```

Weibull and lognormal AFT

Accelerated Failure Time (AFT) models:

$$\log T = \mu + \beta^\top X + \sigma \varepsilon$$

- Weibull and lognormal are two common parametric choices
- Coefficients exponentiated \Rightarrow time ratios, not hazard ratios
- Time ratio > 1 : longer survival (stretching the time scale)
- Time ratio < 1 : shorter survival (compressing the time scale)

We fit

- Fit Weibull and lognormal AFT models
- Compare their AIC and interpret covariate time ratios

Weibull and Lognormal AFT Models



Fit parametric AFT models

```
## Parametric models: Weibull and Lognormal
weibullAFT <- survreg(
  Surv(time, status) ~ temp+load+
    material+maint,
  data = dat,
  dist = "weibull"
)
lognormalAFT <- survreg(
  Surv(time, status) ~ temp+load+
    material+maint,
  data = dat,
  dist = "lognormal"
)
summary(weibullAFT)
summary(lognormalAFT)
AIC(weibullAFT, lognormalAFT)
```

```
Call:
survreg(formula = Surv(time, status) ~ temp + load + material +
  maint, data = dat, dist = "weibull")
      Value Std. Error     z      p
(Intercept) 6.8823   0.9072 7.59 3.3e-14
temp        0.0142   0.0252  0.56  0.5738
load       -1.7517   0.6287 -2.79  0.0053
materialB   0.3730   0.2032  1.84  0.0663
maint1      0.6750   0.2190  3.08  0.0021
Log(scale)  -0.0645   0.0989 -0.65  0.5143

Scale= 0.938

Weibull distribution
Loglik(model)= -611.6  Loglik(intercept only)= -623
Chisq= 22.8 on 4 degrees of freedom, p= 0.00014
Number of Newton-Raphson Iterations: 8
n= 500

> summary(lognormalAFT)

Call:
survreg(formula = Surv(time, status) ~ temp + load + material +
  maint, data = dat, dist = "lognormal")
      Value Std. Error     z      p
(Intercept) 7.5431   1.0421 7.24 4.5e-13
temp        -0.0018   0.0298 -0.06  0.95181
load       -1.8864   0.7292 -2.59  0.00968
materialB   0.2843   0.2421  1.17 0.24031
maint1      0.8388   0.2514  3.34  0.00085
Log(scale)  0.6139   0.0867  7.08 1.4e-12

Scale= 1.85

Log Normal distribution
Loglik(model)= -614.5  Loglik(intercept only)= -624.7
Chisq= 20.39 on 4 degrees of freedom, p= 0.00042
Number of Newton-Raphson Iterations: 4
n= 500

> ## compare AIC
> AIC(weibullAFT, lognormalAFT)
      df      AIC
weibullAFT 6 1235.151
lognormalAFT 6 1240.910
```



Fit parametric AFT models

```
## convert AFT coefficients to time
## ratios (exp on log-time scale)
weibullAFT_timeRatio <- broom::tidy(weibullAFT) %>%
  filter(term != "(Intercept)") %>%
  mutate(
    timeRatio = exp(estimate),
    lower      = exp(estimate - 1.96 * std.error),
    upper      = exp(estimate + 1.96 * std.error)
  )
weibullAFT_timeRatio
```


term	estimate	std.error	statistic	p.value	timeRatio	lower	upper
<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1 temp	0.0142	0.0252	0.562	0.574	1.01	0.965	1.07
2 load	-1.75	0.629	-2.79	0.00533	0.173	0.0506	0.595
3 materialB	0.373	0.203	1.84	0.0663	1.45	0.975	2.16
4 maint1	0.675	0.219	3.08	0.00205	1.96	1.28	3.02
5 Log(scale)	-0.0645	0.0989	-0.652	0.514	0.938	0.772	1.14

Time Ratios from Weibull AFT



Visualizing time ratios

```
## plot time ratios with CI
ggplot(weibullAFT_timeRatio,
       aes(x = term, y = timeRatio,
            ymin = lower, ymax = upper)) +
  geom_pointrange() +
  coord_flip() +
  theme_minimal() +
  geom_hline(yintercept = 1,
             linetype = "dashed") +
  labs(
    x = "Covariate",
    y = "Time ratio (Weibull AFT)",
    title = "AFT (Time Ratios)"
)
```





Checking parametric fit

Compare...

- Kaplan–Meier curve for a subset of pumps
- Weibull AFT fitted survival curve for the same profile

Questions...


- Does the Weibull curve track the KM reasonably well?
- Where does it deviate (early / late follow-up)?

KM vs Weibull AFT for a Profile



Profile and curves

```
## Weibull and KM visual comparison
## we will look at the profile associated with
## Plant P1, material B, maint=1, avg temp/load
profile <- data.frame(temp = mean(dat$temp),
  load    = mean(dat$load),
  material = factor("B", levels = levels(dat$material)),
  maint   = factor(1, levels = levels(dat$maint)))
## Predicted median survival time (months)
predict(weibullAFT, newdata = profile, type = "quantile", p = 0.5)
## Weibull survival curve from survureg params
t_grid <- seq(1, max(dat$time), length.out = 200)
xb     <- predict(weibullAFT, newdata = profile, type = "lp")
sigma  <- weibullAFT$scale
survivalWeibull <- exp(-(t_grid*exp(-xb))^(1/sigma))
weibullDF <- data.frame(time = t_grid, S = survivalWeibull)
kmSUBSET <- survfit(Surv(time, status) ~ 1,
  data = dat %>% filter(material == "B", maint == 1))
gKM <- ggsurvplot(kmSUBSET, conf.int   = TRUE,
  risk.table = FALSE, ggtheme = theme_minimal(),
  xlab = "Time since installation (months)",
  ylab = "Survival probability",
  title = "KM vs Weibull AFT (Material B, Maintained)")
gKM$plot + geom_line(data = weibullDF,
  aes(x = time, y = S), linetype = "dashed") +
  annotate("text", x = max(dat$time)*0.6,
  y = 0.25, label = "Dashed is Weibull AFT")
```



Approximate changing hazards

Idea: model the hazard as constant within time intervals, but allow it to change between intervals.

- Split follow-up into intervals, e.g. $[0, 40), [40, 80), [80, \infty)$ months
- Within each interval: exponential model with its own rate
- Implement as a Poisson regression with $\log(\text{time})$ offset

This is a flexible semi-parametric approach that still yields interpretable rate ratios by interval.

Piecewise Exponential via Poisson Regression



Splitting time and fitting model

```
## Piecewise Exponential Model
cuts <- c(0, 40, 80, Inf)
datSPLIT <- survSplit(Surv(time, status) ~.,
  data = dat, cut = cuts[-c(1, length(cuts))],
  episode = "interval")
## interval factor labels
datSPLIT$interval <- factor(datSPLIT$interval,
  labels = c("0-40", "40-80", "80+"))
## poisson regression on event counts
## with log(time) offset
poissonWeightedGLM <- glm(status ~ interval+temp+
  load+material+maint, family = poisson(),
  offset = log(time), data = datSPLIT)
summary(poissonWeightedGLM)
## coefficients conversion
pweRR <- broom::tidy(poissonWeightedGLM) %>%
  mutate(rate_ratio = exp(estimate),
  lower = exp(estimate-1.96*std.error),
  upper = exp(estimate+1.96*std.error))
pweRR
```

```
Call:
glm(formula = status ~ interval + temp + load + material + maint,
  family = poisson(), data = datSPLIT, offset = log(time))

Coefficients:
Estimate Std. Error z value Pr(>|z|)
(Intercept) -7.14409  0.94012 -7.599 2.98e-14 ***
interval40-80 -0.66334  0.22773 -2.913 0.00358 **
interval80+   -1.35254  0.33826 -3.999 6.37e-05 ***
temp        -0.01633  0.02688 -0.608 0.54343
load         1.85783  0.64821  2.866 0.00416 **
materialB    -0.38419  0.21403 -1.795 0.07264 .
maint1       -0.69620  0.22498 -3.094 0.00197 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 580.97 on 1158 degrees of freedom
Residual deviance: 536.12 on 1152 degrees of freedom
AIC: 728.12

Number of Fisher Scoring iterations: 6

> pweRR
# A tibble: 7 × 8
  term      estimate std.error statistic p.value rate_ratio  lower  upper
  <chr>      <dbl>     <dbl>     <dbl>     <dbl>      <dbl>     <dbl>
1 (Intercept) -7.14      0.940    -7.60  2.98e-14  0.000790  0.000125  0.00498
2 interval40-80 -0.663     0.228    -2.91   3.58e- 3  0.515     0.330    0.805
3 interval80+   -1.35      0.338    -4.00  6.37e- 5  0.259     0.133    0.502
4 temp        -0.0163    0.0269    -0.608 5.43e- 1  0.984     0.933    1.04
5 load         1.86      0.648     2.87  4.16e- 3  6.41      1.80     22.8
6 materialB    -0.384     0.214    -1.80  7.26e- 2  0.681     0.448    1.04
7 maint1       -0.696     0.225    -3.09  1.97e- 3  0.498     0.321    0.775
```


Piecewise Exponential via Poisson Regression



Splitting time and fitting model

```
## interval-specific baseline rates
intervalEFFECTS <- pweRR %>%
  filter(grepl("^interval", term))

ggplot(intervalEFFECTS,
       aes(x = term, y = rate_ratio,
            ymin = lower, ymax = upper)) +
  geom_pointrange() + coord_flip() +
  theme_minimal() +
  labs(
    x = "Interval",
    y = "Rate ratio versus reference",
    title = "Failure Rate by Time Interval"
)
```



Which model is best?

Compare...

- AIC: penalized likelihood; lower is better (but scale differs across model classes)
- Harrell's C-index: rank-based discrimination; 0.5 = random, 1.0 = perfect
- Prediction error / Brier score: absolute prediction error over time

Notes...

- Compare Cox (with/without frailty) to Weibull and lognormal AFT
- Use `pec` to obtain prediction error curves for Cox-type models

AIC, C-index, and Prediction Error



Model comparison code

```
## Which model predicts best?
AIC(coxBASIC, coxFRAILTY, weibullAFT, lognormalAFT)

## Predictive accuracy with Harrell's C-index
## C-index measures rank-order predictive accuracy
cINDEX <- function(fit, data) {
  s <- Surv(data$time, data$status)
  lp <- predict(fit, newdata = data, type = "lp")
  conc <- concordance(s ~ lp)$concordance
  return(conc)
}
c_basic <- cINDEX(coxBASIC, dat)
c_frail <- cINDEX(coxFRAILTY, dat)
c_weibull <- cINDEX(weibullAFT, dat)
c_logn <- cINDEX(lognormalAFT, dat)

data.frame(
  model = c("Cox (no frailty)",
            "Cox (frailty)", "Weibull AFT", "Lognormal AFT"),
  cINDEX = c(c_basic, c_frail, c_weibull, c_logn)
)
```

	df	AIC
coxBASIC	4.000000	1033.967
coxFRAILTY	5.304768	1028.984
weibullAFT	6.000000	1235.151
lognormalAFT	6.000000	1240.910


	model	cINDEX
1	Cox (no frailty)	0.3621266
2	Cox (frailty)	0.3423481
3	Weibull AFT	0.6376427
4	Lognormal AFT	0.6407565

AIC, C-index, and Prediction Error



Model comparison code

```
## Brier score absolute prediction error at time t.
## still figuring out how to do this
## for the other two models
models <- list(
  coxBASIC = coxBASIC,
  coxFRAILTY = coxFRAILTY
)
## prediction error curves
pecFIT <- pec(
  object = models,
  formula = Surv(time, status) ~ 1,
  data    = dat,
  times   = quantile(dat$time,
    probs = seq(0.1, 0.9, by = 0.1)))
plot(pecFIT,
  xlab = "Time (months)",
  ylab = "Prediction error",
  legend = TRUE,
  main = "Prediction Error Curves")
```





Additional Questions?
Book an Appointment!



Next Workshop:

Nonparametric Models
→ December 2, 11:00 AM

Thank You!

Questions?

Workshop Materials:

<https://github.com/csc-ubc-okanagan/ubco-csc-modeling-workshop>

Contact:

jesse.ghashti@ubc.ca